



Logical Schema Design that Quantifies Update Inefficiency and Join Efficiency

Sebastian Link
 s.link@auckland.ac.nz
 The University of Auckland
 Auckland, New Zealand

Ziheng Wei
 z.wei@auckland.ac.nz
 The University of Auckland
 Auckland, New Zealand

ABSTRACT

The goal of classical normalization is to maintain data consistency under updates, with a minimum level of effort. Given functional dependencies (FDs) alone, this goal is only achievable in the special case an FD-preserving Boyce-Codd Normal Form (BCNF) decomposition exists. As we show, in all other cases the level of effort can be neither controlled nor quantified. In response, we establish the ℓ -Bounded Cardinality Normal Form, parameterized by a positive integer ℓ . For every ℓ , the normal form condition requires from every instance that every value combination over the left-hand side of every non-trivial FD does not occur in more than ℓ tuples. BCNF is captured when $\ell = 1$. We demonstrate that schemata in this normal form characterize the instances that are i) free from level ℓ data redundancy and update inefficiency, and ii) permit level ℓ join efficiency. We establish algorithms that compute schemata in ℓ -Bounded Cardinality Normal Form for the smallest level ℓ attainable across all FD-preserving decompositions. Additional algorithms i) attain even smaller levels of effort based on the loss of some FDs, and ii) decompose schemata based on prioritized FDs that cause high levels of effort. Our framework informs de-normalization already during logical design. In particular, level ℓ quantifies both the incremental maintenance and join support of materialized views. Experiments with synthetic and real-world data illustrate which properties the schemata have that result from our algorithms, and how these properties predict the performance of update and query operations on instances over the schemata, without and with materialized views.

CCS CONCEPTS

- Information systems → Database design and models; Relational database model; Integrity checking.

KEYWORDS

Cardinality constraint; Data redundancy; Functional dependency; Join; Normal form; Normalization; Update

ACM Reference Format:

Sebastian Link and Ziheng Wei. 2021. Logical Schema Design that Quantifies Update Inefficiency and Join Efficiency. In *Proceedings of the 2021*

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SIGMOD '21, June 20–25, 2021, Virtual Event, China

© 2021 Copyright held by the owner/author(s). Publication rights licensed to ACM. ACM ISBN 978-1-4503-8343-1/21/06...\$15.00
<https://doi.org/10.1145/3448016.3459238>

International Conference on Management of Data (SIGMOD '21), June 20–25, 2021, Virtual Event, China. ACM, New York, NY, USA, 13 pages. <https://doi.org/10.1145/3448016.3459238>

1 INTRODUCTION

Schema design aims at finding a layout of the data that facilitates the efficient processing of common queries and updates. The problem is challenging as data redundancy typically causes update inefficiency but promotes join efficiency. So far, the challenge has been addressed by performing normalization during logical schema design to achieve update efficiency, followed by de-normalization during physical design to boost query efficiency. In particular, de-normalization is done only after the database is operational and patterns of data access on normalized databases emerge.

The goal of classical normalization is to maintain data consistency under updates, with a minimum level of effort. Normalization aims at eliminating any occurrence of redundant data values in any future database instance. This is attempted by structurally transforming functional dependencies (FDs), that cause data redundancy, into keys, that prohibit data redundancy. The well-known Boyce-Codd Normal Form (BCNF) requires the left-hand side of every non-trivial FD to be a key. Hence, data redundancy can never be caused by FDs transformed into keys. However, while every schema can be decomposed into BCNF, FDs may be lost during this process and still cause data redundancy [5]. A more liberal condition is given by the Third Normal Form (3NF) where the left-hand side of every non-trivial FD must be a key or every attribute on the right-hand side must be prime (that is, to be part of some minimal key). 3NF synthesis can transform every schema into 3NF without losing any FD, but some FDs still cause data redundancy as they were not transformed into keys. Hence, classical normalization can only measure its success when no effort is required at all to maintain data consistency. This is only possible when an FD-preserving BCNF decomposition exists. In all other cases, the level of effort can be neither controlled nor even measured.

As running example consider the event management schema HAP with each record representing an *E(vent)* at a *V(enue)* and *T(ime)* with some *C(ompany)* in charge. HAP uses FDs to model business rules. The FD $E \rightarrow C$ says that only one company is in charge of every event, $V \rightarrow C$ says there cannot be different companies in charge at the same venue, $VT \rightarrow E$ says that no different events happen at the same venue at the same time, $ET \rightarrow V$ says that no event takes place at different venues at the same time, and $CT \rightarrow V$ expresses that no company is in charge of different venues at the same time. The minimal keys are ET , VT and CT . Hence, every attribute of HAP is prime. While HAP is in 3NF it has

Table 1: Schema HAP in 3NF and its decomposition \mathcal{D}_1 together with instances (redundant data values in bold)

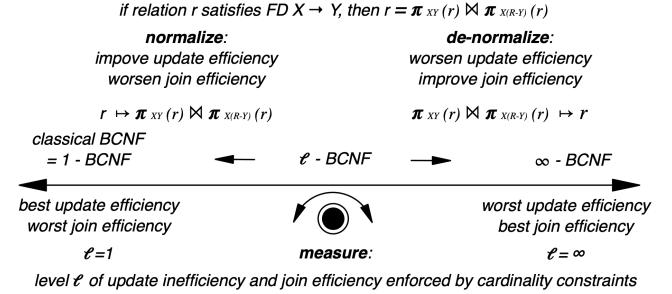
(a) Instance r over HAP				(b) Decomposition $\mathcal{D}_1 = \{R_2, R_3, R_4\}$ of HAP with projected instance $\{r_i = \pi_{R_i}(r)\}_{i=2}^4$							
HAP				R_2		R_3		R_4			
<i>Event</i>	<i>Venue</i>	<i>Company</i>	<i>Time</i>	<i>Venue</i>	<i>Company</i>	<i>Event</i>	<i>Venue</i>	<i>Time</i>	<i>Event</i>	<i>Company</i>	<i>Time</i>
Party	v_1	Kilo	t_1	v_1	Kilo	Party	v_1	t_1	Party	Kilo	t_1
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
Party	v_{ℓ_1}	Kilo	t_{ℓ_1}	v_{ℓ_1}	Kilo	Party	v_{ℓ_1}	t_{ℓ_1}	Party	Kilo	t_{ℓ_1}
e_1	Dome	Mega	t'_1	Dome	Mega	e_1	Dome	t'_1	e_1	Mega	t'_1
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
e_{ℓ_2}	Dome	Mega	t'_{ℓ_2}			e_{ℓ_2}	Dome	t'_{ℓ_2}	e_{ℓ_2}	Mega	t'_{ℓ_2}

no FD-preserving BCNF decomposition. Given FD-preservation, classical normalization cannot achieve more.

However, we cannot measure the level of effort to maintain data consistency for HAP. In instance r of Table 1a each of the $\ell_1 > 1$ occurrences of C -value *Kilo* is redundant due to the FD $E \rightarrow C$, and each of the $\ell_2 > 1$ occurrences of C -value *Mega* is redundant due to $V \rightarrow C$. Recall that a value occurrence is redundant whenever every update of the occurrence to a different value results in a violation of the FD. We refer to the number of different tuples in which a given data value can occur redundantly as *the level of data redundancy*. Clearly, the level of data redundancy on HAP is unbounded. This is true for every schema that is in 3NF but not in BCNF. As we will show, every FD lost during BCNF decomposition still causes an unbounded level of data redundancy. Hence, the number of data values that need updating is a priori unbounded. Changing one occurrence of *Kilo* means that a total of ℓ_1 occurrences of *Kilo* need updating to ensure data consistency. We refer to the total number of occurrences that require an update to achieve consistency as the *level of update inefficiency*. So, for FD-preserving BCNF decompositions the level of effort for data consistency is at optimum 1, while it is unbounded in all other cases.

In practice, however, ℓ_1 and ℓ_2 represent upper bounds of *cardinality constraints* (CCs). They constitute a different class of business rules than FDs. For a positive integer ℓ , the CC $\text{card}(X) \leq \ell$ says that every instance can have up to ℓ different records with matching values on all the attributes in X . For $\ell = 1$, X is a key. For $\ell = \infty$, no bound has been specified. For example, the CC $\text{card}(E) \leq \ell_1$ with $\ell_1 = 1,000$ (1k) expresses that every event can have up to 1k different combinations of venues and times. Similarly, the CC $\text{card}(V) \leq \ell_2$ with $\ell_2 = 1,000,000$ (1m) expresses that every venue can have up to 1m different combinations of events and times. CCs inform schema design beyond classical normalization. Given CCs, the 3NF schema HAP admits level ℓ_2 data redundancy and update inefficiency. Without CCs no integer bounds are specified, and every non-trivial or lost FD causes an unbounded level of data redundancy and update inefficiency, unless its left-hand side is a key. Hence, we propose to include CCs in schema design.

With CCs we can quantify the level of effort required to maintain data consistency under updates. For example, we may ask which FD-preserving decompositions of HAP minimize level ℓ . Applying one of our new algorithms results in schema $\mathcal{D}_1 = \{(R_2, \Sigma_2), (R_3, \Sigma_3), (R_4, \Sigma_4)\}$ in Table 1b with

**Figure 1: Level ℓ of Effort Required For Data Consistency, and its Impact on Update and Join Efficiency**

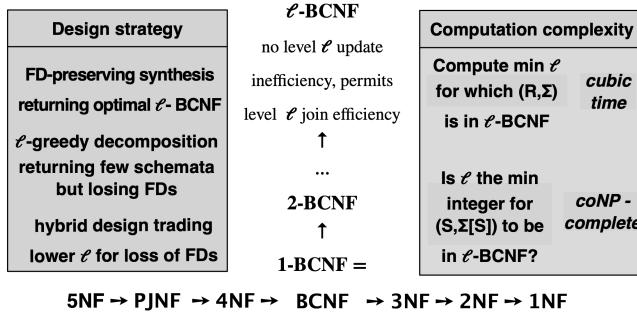
- $R_2 = CV$ and $\Sigma_2 = \{V \rightarrow C\}$,
- $R_3 = ETV$ and $\Sigma_3 = \{TV \rightarrow E, ET \rightarrow V\}$, and
- $R_4 = CET$ and $\Sigma_4 = \{CT \rightarrow E, E \rightarrow C\}$.

The schemata (R_2, Σ_2) and (R_3, Σ_3) are both in BCNF and cannot exhibit any redundant data value. However, R_4 is in 3NF and still exhibits an $\ell_1 = 1k$ level of data redundancy and update inefficiency caused by $E \rightarrow C$. Indeed, $\ell_1 = 1k$ is the optimum level achievable by any FD-preserving decomposition of HAP. The given FD $CT \rightarrow V$ is preserved by \mathcal{D}_1 as it is implied by $CT \rightarrow E$ and $ET \rightarrow V$. The decomposition $\mathcal{D}_2 = \{(R_1, \Sigma_1), (R_3, \Sigma_3), (R_5, \Sigma_5)\}$ with

- $R_1 = EC$ and $\Sigma_1 = \{E \rightarrow C\}$, and
- $R_5 = CTV$ and $\Sigma_5 = \{CT \rightarrow V, V \rightarrow C\}$

is in 3NF and achieves level $\ell_2 = 1m$ data redundancy.

CCs were already introduced in Chen's seminal ER paper [9] and used by UML, XML, and OWL. Surprisingly, they have not been used for normalization. We also observe that CCs quantify *the level of join efficiency* for a schema. We define the latter as the maximum number of values an FD can join with some redundant value in any instance of the schema. Instance r_4 over R_4 in Table 1b satisfies $E \rightarrow C$. Hence, r_4 is equal to the join $\pi_{EC}(r_4) \bowtie \pi_{ET}(r_4)$. Indeed, the redundant data C -value *Kilo* can be joined with the ℓ_1 T -values t_1, \dots, t_{ℓ_1} via the E -value *Party*. Without CCs, the join strength is unbounded, so the classical theory does not provide insight into the join efficiency of a schema. CCs inform the selection of materialized views, even before the database is operational. Here the level ℓ of a materialized view quantifies both its effort required for incremental

Figure 2: Achievements of ℓ -BCNF

maintenance and its support for joins. Defining HAP as view over D_1 , its level of maintenance and join efficiency is $\ell_2 = 1m$.

Contributions. Our two main contributions are:

- We establish the first framework for logical schema normalization that quantifies the effort required to achieve data consistency during updates.
- Our framework also quantifies the join efficiency of schemata, which means it informs de-normalization already during logical design.

Figure 1 illustrates our main ideas, in particular that ℓ quantifies both the update effort and join capability of schemata. Figure 2 illustrates our technical results, which are:

- (1) We relax the classical BCNF condition by permitting every value combination over the left-hand side of every non-trivial FD to occur in up to ℓ tuples of every instance. Hence, we obtain the infinite hierarchy of ℓ -Bounded Cardinality Normal Forms (ℓ -BCNF), with classical BCNF captured for $\ell = 1$. Therefore, the minimum ℓ for which ℓ -BCNF is attainable measures the effort required to achieve data consistency.
- (2) We show for every ℓ that schemata in ℓ -BCNF characterize instances that are i) free from level ℓ data redundancy and update inefficiency, and ii) permit a level ℓ of join efficiency.
- (3) We establish an algorithm that computes schemata in ℓ -BCNF for the minimum level ℓ attainable across FD-preserving decompositions. Another algorithm prioritizes FDs for lossless decompositions based on the level of data redundancy they cause. This algorithm produces few output schemata, but FDs may be lost. Combining both algorithms to a hybrid strategy, we further reduce the minimum level attainable from FD-preserving decompositions by losing FDs.
- (4) Level ℓ informs the selection of materialized views as it captures both the incremental effort to maintain the views under updates, and the support for join queries by the views.
- (5) Experiments with synthetic and real-world data illustrate which properties the schemata have that result from our algorithms, and how these properties predict the physical performance of update and query operations on instances over the schemata, without and with materialized views.

Organization. We recall preliminaries in Section 2. In Section 3 we introduce our family of ℓ -BCNF and its computational properties. In Section 4 we establish which properties schemata in ℓ -BCNF exhibit in terms of updates and joins. We discuss three design algorithms in

Section 5. Experimental results are presented in Section 6. Related work is discussed in Section 7. We conclude in Section 8. A technical report contains details and links to data sets [36].

2 PRELIMINARIES

Classical normalization for FDs is part of most textbooks for relational databases. Table 2 recalls the definition and notation we will use here. In addition, we explicitly recall less familiar concepts around CCs. For a detailed background on classical normalization we refer the reader to some classical literature [3, 4, 38].

Let $\mathbb{N}_{\geq 1}^{\infty}$ denote the positive integers and ∞ . A *cardinality constraint* (CC) over relation schema R is an expression $card(X) \leq \ell$ where $X \subseteq R$ and $\ell \in \mathbb{N}_{\geq 1}^{\infty}$. A relation r over R satisfies $card(X) \leq \ell$ whenever there are no more than ℓ different tuples in r that all have matching values on all the attributes in X , that is,

$$\forall t_1, \dots, t_{\ell+1} \in r \quad (t_1(X) = \dots = t_{\ell+1}(X) \Rightarrow \exists i, j \in \{1, \dots, \ell+1\} (t_i = t_j)).$$

For example, the relation in Table 1a satisfies $card(EC) \leq \ell_1 = 1k$ but violates $card(EC) \leq 999$ due to the 1k different tuples that have *E*-value *Party* and *C*-value *Kilo*.

Table 3 shows Armstrong's axioms as set $\mathfrak{A} = \{\mathcal{R}, \mathcal{E}, \mathcal{T}\}$ which forms an axiomatization for FDs alone [2], the set $\{\mathcal{S}, \mathcal{U}, \mathcal{L}, \mathcal{A}\}$ which forms an axiomatization for CCs alone [21], and the set \mathfrak{L} of all rules in Table 3 which forms an axiomatization for CCs and FDs together [21]. In our example, we may apply rule \mathcal{E} to $E \rightarrow C$ to infer $E \rightarrow EC$. We then apply \mathcal{P} to $E \rightarrow EC$ and $card(EC) \leq \ell_1$ to infer $card(E) \leq \ell_1$.

Axiomatizations target efficient algorithms. An FD $X \rightarrow Y$ is implied by an FD set Σ if and only if Y is a subset of the attribute set closure X_{Σ}^{+} . This has resulted in a linear-time algorithm for FD implication [3]. Though more expressive than CCs and FDs in isolation, their combined implication can be reduced to that for FDs alone by translating any set Σ of CCs and FDs into the FD set

$$\begin{aligned} \Sigma[FD] = & \{X \rightarrow Y \mid X \rightarrow Y \in \Sigma\} \cup \\ & \{X \rightarrow R \mid card(X) \leq 1 \in \Sigma\} [21]. \end{aligned}$$

Essentially, an FD $X \rightarrow Y$ is implied by the CC/FD set Σ if and only if it is implied by the FD set $\Sigma[FD]$, and a CC $card(X) \leq \ell$ is implied by the CC/FD set Σ if and only if there is some $card(Y) \leq \ell' \in \Sigma \cup \{card(R) \leq 1\}$ such that $X \rightarrow Y$ is implied by $\Sigma[FD]$ and $\ell' \leq \ell$ [21].

For example, if Σ contains $E \rightarrow C$ and $card(EC) \leq 1k$, then $card(E) \leq 2k$ is implied by Σ since $E_{\Sigma[FD]}^{+} = EC$ and $EC \subseteq E_{\Sigma[FD]}^{+}$ for the CC $card(EC) \leq 1k \in \Sigma$ and $\ell' = 1k \leq 2k = \ell$.

So, deciding implication for CCs and FDs combined is in time $O(|\Sigma| \times ||\Sigma||)$ where $|\Sigma|$ is the number of elements in Σ and $||\Sigma||$ the total number of attribute occurrences in Σ [21].

3 THE FAMILY OF ℓ -BCNF

We establish the family of ℓ -Bounded Cardinality Normal Forms and some computational properties. Its levels of update inefficiency and join efficiency are studied in Section 4.

A schema is in BCNF whenever the left-hand side (LHS) X of every non-trivial FD $X \rightarrow Y$ is a key. Since X is a key whenever $card(X) \leq 1$ holds, CCs provide a convenient mechanism to relax the BCNF condition as follows.

Table 2: Concepts and their Definition from the Relational Normalization Framework

<i>Concept</i>	<i>Definition</i>
relation schema R	finite set R of attributes A , each with a domain $\text{dom}(A)$ of possible values
database schema \mathcal{D}	finite set $\mathcal{D} = \{R_1, \dots, R_n\}$ of relation schemata
tuple t over R	function $t : R \rightarrow \bigcup_{A \in R} \text{dom}(A)$ with $t(A) \in \text{dom}(A)$
relation r over R	finite set r of tuples over R
projection of relation r on X	$r(X) = \{t(X) \mid t \in r\}$ where $t(X)$ is the projection of tuple t on attribute set X
r satisfies FD $X \rightarrow Y$ ($\models_r X \rightarrow Y$)	for all tuples $t, t' \in r$, if $t(X) = t'(X)$, then $t(Y) = t'(Y)$
r satisfies FD set Σ ($\models_r \Sigma$)	for all $\sigma \in \Sigma$, $\models_r \sigma$
Σ implies φ over R ($\Sigma \models \varphi$)	for all relations r over relation schema R : if $\models_r \Sigma$, then $\models_r \varphi$
inference of φ from Σ ($\Sigma \vdash_{\mathfrak{R}} \varphi$)	φ can be inferred from Σ by applying rules from \mathfrak{R}
semantic closure of FD set Σ	$\Sigma^* = \{\varphi \mid \Sigma \models \varphi\}$ (all the FDs implied by Σ)
syntactic closure of FD set Σ	$\Sigma_{\mathfrak{R}}^+ = \{\varphi \mid \Sigma \vdash_{\mathfrak{R}} \varphi\}$ (all the FDs that can be inferred from Σ by applying rules in \mathfrak{R})
attribute set closure of X under Σ	$X_{\Sigma}^+ = \{A \in R \mid X \rightarrow A \in \Sigma_{\mathfrak{R}}^+\}$ (all the attributes functionally determined by attribute set X)
axiomatization, denoted by \mathfrak{R}	set \mathfrak{R} of rules that is sound: $\Sigma^* \subseteq \Sigma_{\mathfrak{R}}^+$ and complete: $\Sigma_{\mathfrak{R}}^+ \subseteq \Sigma^*$
cover of FD set Σ	an FD set Θ is a cover of Σ if and only if $\Theta^* = \Sigma^*$ (Θ implies the same FDs as Σ does)
non-redundant (L-reduced) FD set Σ	for all $\sigma \in \Sigma$, $\Sigma - \{\sigma\} \not\models \sigma$ (and $\forall X \rightarrow Y \in \Sigma$ there is no proper $Z \subset X$ such that $\Sigma \models Z \rightarrow Y$)
canonical cover Σ	cover Σ that is non-redundant, L-reduced, and left-hand sides are unique
Armstrong axioms, denoted by \mathfrak{A}	rule set with reflexivity axiom \mathcal{R} , extension rule \mathcal{E} , and transitivity rule \mathcal{T} (see Table 3)
(minimal) key for (R, Σ)	(minimal) $X \subseteq R$ such that $\Sigma \models X \rightarrow R$
prime attribute $A \in R$	attribute A contained in some minimal key for (R, Σ)
(R, Σ) in BCNF (3NF)	for every non-trivial FD $X \rightarrow Y \in \Sigma_{\mathfrak{R}}^+$, X is key for (R, Σ) (or every $A \in Y - X$ is prime)
(lossless) decomposition of (R, Σ)	database schema \mathcal{D} with $\bigcup_{S \in \mathcal{D}} S = R$ (and for all R -relations r : if $\models_r \Sigma$, then $r = \bowtie_{S \in \mathcal{D}} r(S)$)
projection $\Sigma[S]$ of FD set Σ to S	$\Sigma[S] = \{X \rightarrow Y \in \Sigma^* \mid X, Y \subseteq S\}$ (all the FDs implied by Σ with attributes in attribute set S)
BCNF (3NF) decomposition of (R, Σ)	decomposition \mathcal{D} of (R, Σ) such that for all $S \in \mathcal{D}$, $(S, \Sigma[S])$ is in BCNF (3NF)
FD-preserving decomposition of (R, Σ)	decomposition \mathcal{D} of (R, Σ) such that $\forall X \rightarrow Y \in \Sigma ((\bigcup_{S \in \mathcal{D}} \Sigma[S]) \models X \rightarrow Y)$

Table 3: Axiomatizations of CCs and FDs

$\frac{\text{card}(R) \leq 1}{(\text{set}, \mathcal{S})}$	$\frac{\text{card}(X) \leq \infty}{(\text{unbounded}, \mathcal{U})}$	
$\frac{\text{card}(X) \leq \ell}{\text{card}(X) \leq \ell + 1}$	$\frac{\text{card}(X) \leq \ell}{\text{card}(XY) \leq \ell}$	
(loosen, \mathcal{L})	(adding, \mathcal{A})	
$\frac{XY \rightarrow Y}{(\text{reflexivity}, \mathcal{R})}$	$\frac{X \rightarrow Y}{X \rightarrow XY}$	$\frac{X \rightarrow Y \quad Y \rightarrow Z}{X \rightarrow Z}$
(extension, \mathcal{E})		(transitivity, \mathcal{T})
$\frac{\text{card}(X) \leq 1}{X \rightarrow Y}$	$\frac{X \rightarrow Y \quad \text{card}(Y) \leq \ell}{\text{card}(X) \leq \ell}$	
(key, \mathcal{K})	(pullback, \mathcal{P})	

Definition 3.1. Let Σ denote a set of CCs and FDs over relation schema R , and let $\ell \in \mathbb{N}_{\geq 1}^\infty$. (R, Σ) is in ℓ -Bounded Cardinality Normal Form (ℓ -BCNF) iff for all $X \rightarrow Y \in \Sigma_{\mathfrak{R}}^+$ where $Y \not\subseteq X$ we have that $\text{card}(X) \leq \ell \in \Sigma_{\mathfrak{R}}^+$.

Example 3.2. For our running example HAP consider the FD set Θ of $E \rightarrow C, V \rightarrow C, VT \rightarrow E, ET \rightarrow V$, and $CT \rightarrow V$. While (HAP, Θ) is in 3NF, it is in ∞ -BCNF as the smallest ℓ for which $\text{card}(E) \leq \ell \in \Theta_{\mathfrak{R}}^+$ is $\ell = \infty$. Intuitively, this matches our understanding that the levels of data redundancy, update inefficiency, and join efficiency are unbounded. Consider now (HAP, Σ) where Σ consists of the FDs in Θ

and the CCs $\text{card}(E) \leq 1k$ and $\text{card}(V) \leq 1m$. Since VT, ET , and CT are (minimal) keys, we obtain the CCs $\text{card}(VT) \leq 1, \text{card}(ET) \leq 1$, and $\text{card}(CT) \leq 1$. Hence, every non-trivial $X \rightarrow Y \in \Sigma$ satisfies $\text{card}(X) \leq 1m \in \Sigma_{\mathfrak{R}}^+$. The latter condition is equivalent to being in ℓ -BCNF (Theorem 3.5), so (HAP, Σ) is in $1m$ -BCNF. (HAP, Σ) is not in any ℓ -BCNF for $\ell < 1m$ since we have the FD $V \rightarrow C \in \Sigma$ but the smallest ℓ_2 with $\text{card}(V) \leq \ell_2 \in \Sigma_{\mathfrak{R}}^+$ is $\ell_2 = 1m$.

Our definition of ℓ -BCNF has the desirable property to be invariant under covers, as stated in the next result.

THEOREM 3.3. Let Σ and Θ denote two CC/FD sets over R that are covers of one another. For all $\ell \in \mathbb{N}_{\geq 1}^\infty$, (R, Σ) is in ℓ -BCNF iff (R, Θ) is in ℓ -BCNF.

Theorem 3.3 means we need not worry how we represent application semantics as integrity constraints. For example, let Σ' denote the FDs in Θ together with the CCs $\text{card}(EC) \leq 1k$ and $\text{card}(VC) \leq 1m$. Then Σ and Σ' are covers of one another, as easily seen from the adding rule \mathcal{A} , the extension rule \mathcal{E} and the pullback rule \mathcal{P} . Hence, (HAP, Σ') is in 1m-BCNF but not in ℓ -BCNF for any $\ell < 1m$.

Our definition gives rise to a family of syntactic normal forms that exhibit a strict hierarchy. At the bottom is 1-BCNF, which is equivalent to the classical Boyce-Codd normal form.

THEOREM 3.4. For every schema (R, Σ) and every positive integer ℓ we have: If (R, Σ) is in ℓ -BCNF, then (R, Σ) is also in $\ell + 1$ -BCNF. However, for every positive integer ℓ there are schemata (R, Σ) in $\ell + 1$ -BCNF that are not in ℓ -BCNF.

Algorithm 1 Strongest ℓ -Bounded Cardinality Normal Form

Require: (R, Σ) with CC/FD set Σ over schema R

Ensure: Minimum integer ℓ such that (R, Σ) is in ℓ -BCNF

```

1: if  $\Sigma_{FD} = \emptyset$  then
2:   return 1
3:  $\ell \leftarrow \infty$ 
4:  $\Sigma_{card} \leftarrow \Sigma_{card} \cup \{card(R) \leq 1\}$ 
5: for all  $X \rightarrow Y \in \Sigma_{FD}$  such that  $Y \not\subseteq X$  do
6:   if  $\ell_X$  has not been defined before then
7:      $\ell_X \leftarrow \infty$ 
8:   for all  $card(Y) \leq \ell' \in \Sigma_{card}$  with  $\ell' < \ell_X$  do
9:     if  $Y \subseteq X_{\Sigma[FD]}^+$  then
10:       $\ell_X \leftarrow \ell'$ 
11:   if  $\ell_X \leftarrow \infty$  then
12:     return  $\infty$ 
13: return  $\max_X \{\ell_X\}$ 
```

For example, (HAP, Σ) is in ℓ -BCNF if and only if $\ell \geq 1m$. Theorem 3.4 establishes the first infinite hierarchy of normal forms for relational schema design. It is orthogonal to known normal forms such as 3NF, BCNF, 4NF, etc.

Combining CCs and FDs raises expressiveness without adding (much) computational complexity. As a generalization of BCNF, known hardness results apply to ℓ -BCNF.

3.1 Efficient Decidability Locally

Firstly, we show that deciding whether for a given schema (R, Σ) and for a given positive integer ℓ , (R, Σ) is in ℓ -BCNF can be done in cubic time in the input. This is a consequence of showing that it suffices to check the FDs in Σ to validate whether (R, Σ) is in ℓ -BCNF.

THEOREM 3.5. *For all (R, Σ) and $\ell \in \mathbb{N}_{\geq 1}^\infty$, (R, Σ) is in ℓ -BCNF iff for all $X \rightarrow Y \in \Sigma$ where $Y \not\subseteq X$, we have that $card(X) \leq \ell \in \Sigma_{\Sigma}^+$. We can decide in $O(|\Sigma|^2 \times |\Sigma|)$ time whether for a given schema (R, Σ) and a given $\ell \in \mathbb{N}_{\geq 1}^\infty$, (R, Σ) is in ℓ -BCNF.*

3.2 Computing the strongest local normal form

Our new setting motivates the problem of computing the smallest positive integer ℓ for which a given schema (R, Σ) is in ℓ -BCNF. For the set Σ of CCs and FDs over relation schema R , we use Σ_{FD} to denote the set of FDs in Σ , and Σ_{card} the set of CCs in Σ . Algorithm 1 computes the smallest ℓ for which (R, Σ) is in ℓ -BCNF. The steps are as follows.

If there are no FDs in the input (line 1), the output level will be 1 (line 2). Otherwise, we start with the worst possible level ∞ (line 3). We add the CC $card(R) \leq 1$ to the set of given CCs, since R is always a key. For each non-trivial input FD with LHS X (line 5) we determine the smallest positive integer ℓ_X by which X is bound according to the input (line 5-12). Initially, ℓ_X is ∞ (line 6-7), and the current ℓ_X can be replaced by a smaller ℓ' (line 10) whenever some CC $card(Y) \leq \ell' \in \Sigma_{card}$ can be found (line 8) such that the FD $X \rightarrow Y$ is implied (line 9). Indeed, in that case, the CC $card(X) \leq \ell'$ is also implied by Σ (see the pullback rule \mathcal{P} from Table 3). We terminate with ∞ as soon as we find that $\ell_X = \infty$ for

some X (line 11-12). Otherwise, we return the maximum ℓ_X among those computed (line 13).

Algorithm 1 computes the strongest ℓ -BCNF in the sense that ℓ is the smallest positive integer with that property. Essentially, for each FD $X \rightarrow Y$ in the input we check for each cardinality constraint in the input whether it implies some smaller bound ℓ_X . Algorithm 1 therefore operates in $O(|\Sigma|^2 \times |\Sigma|)$ time. We thus obtain the following result.

COROLLARY 3.6. *Given a set Σ of CCs and FDs over relation schema R , we can compute in $O(|\Sigma|^2 \times |\Sigma|)$ time the smallest positive integer ℓ such that (R, Σ) is in ℓ -BCNF.*

For (HAP, Σ) , $\ell_E = 1k$, $\ell_V = 1m$, $\ell_{VT} = 1$, $\ell_{ET} = 1$, and $\ell_{CT} = 1$. Since the maximum is $\ell = \ell_2 = 1m$, ℓ_2 is the optimum ℓ for which (HAP, Σ) is in ℓ -BCNF.

3.3 Likely Intractability Globally

While we can efficiently compute the strongest normal form locally, this is unlikely to be efficient globally. Given (R, Σ, ℓ) and a subschema $S \subset R$, it is unlikely we can find a PTIME algorithm deciding if $(S, \Sigma[S])$ is in ℓ -BCNF.

THEOREM 3.7. *Let Σ denote a set of CCs and FDs over $R, S \subset R$, and ℓ a positive integer. Given (R, S, Σ, ℓ) , it is coNP-complete to decide whether $(S, \Sigma[S])$ is in ℓ -BCNF.*

While it is already coNP-complete to decide if a given subschema is in BCNF (1-BCNF), Theorem 3.7 encourages us to view schema normalization from the new perspective of computing the smallest ℓ for which a schema is in ℓ -BCNF.

4 USEFUL PROPERTIES OF ℓ -BCNF

We show that schemata in ℓ -BCNF characterize instances with useful properties for updates and joins. In particular, ℓ quantifies the highest number of values that need updating to achieve data consistency, but also how many data values can be joined with a given redundant value.

4.1 ℓ -redundancy

Intuitively, Vincent [47] defined a single occurrence of a data value as *redundant* whenever every change to this occurrence results in a relation that violates some given constraint. Our idea is to fix some positive integer ℓ , and define a data value as *ℓ -redundant* whenever there are ℓ distinct tuples in which the value occurs and every update to at least one of these ℓ occurrences results in a relation that violates a given constraint. That is, the value of these ℓ occurrences is already uniquely determined by the other values and knowing the constraints are satisfied by the relation.

For example, the value *Kilo* in relation r_4 in Table 1a is 999-but not 1k-redundant ($\ell_1 = 1k$): Concealing between 1 and 999 occurrences of the value *Kilo* in the *C*-column still allows us to determine each of the concealed occurrences as $E \rightarrow C$ must hold and there is at least one tuple left that has a matching *E*-value and *C*-value *Kilo*, as illustrated next.

E	C	T		E	C	T
Party	?	t_1	$E \rightarrow C$	Party	Kilo	t_1
:	:	:	$= \text{Kilo}$:	:	:
Party	?	t_{999}		Party	Kilo	t_{999}
Party	Kilo	t_{1k}		Party	Kilo	t_{1k}

However, concealing all $1k$ occurrences means the tuples could have any value on C (as long as they all match on C).

We define a family of semantic normal forms by stipulating the absence of any relations that feature any ℓ -redundancy. A schema (R, Σ) is in ℓ -Redundancy Free Normal Form (ℓ -RFNF) iff there is no relation r over R that satisfies Σ , there is no $\sigma \in \Sigma$, and there is no attribute $A \in R$ for which there is a data value $v \in r(A)$ that is ℓ -redundant for σ .

(HAP, Σ) from Example 3.2 is in $1m$ -RFNF ($\ell_2 = 1m$), but not in ℓ -RFNF for any $\ell < \ell_2$. Concealing any fewer than the ℓ_2 occurrences of *Mega* will allow us to deduce their value.

More formally, the data value $v \in r(A)$ is ℓ -redundant for $\sigma \in \Sigma$ iff there are ℓ distinct tuples $t_1, \dots, t_\ell \in r$ such that $t_1(A) = \dots = t_\ell(A) = v$, and for every ℓ -transaction $\{t'_1, \dots, t'_\ell\}$ of t_1, \dots, t_ℓ for A , the relation $r' := (r - \{t_1, \dots, t_\ell\}) \cup \{t'_1, \dots, t'_\ell\}$ violates σ .

An ℓ -transaction causes an actual update to at least one of ℓ values. More formally, an ℓ -transaction of t_1, \dots, t_ℓ for A is a set $\{t'_1, \dots, t'_\ell\}$ of tuples over R such that for all $i = 1, \dots, \ell$ and for all $A' \in R - \{A\}$, $t'_i(A') = t_i(A')$, and there is some $j \in \{1, \dots, \ell\}$ such that $t'_j(A) \neq t_j(A)$.

4.2 ℓ -update inefficiency

An ℓ -redundant data value is synonymous with that of an ℓ -update inefficiency in the following sense. No matter how any ℓ occurrences of the ℓ -redundant data value are updated, consistency cannot be achieved since there will still remain some occurrence of the value that could not have been updated. For the table on the right above, no matter how we update 999 occurrences of *Kilo*, as long as we update at least one of them we cannot satisfy $E \rightarrow C$. However, updating all 1000 occurrences of *Kilo* consistently to a new value will satisfy $E \rightarrow C$. Based on the CC $\text{card}(E) \leq 1k$, we require the update of at most 1000 tuples to maintain consistency for $E \rightarrow C$, in any legal relation. Hence, it is justified to say that (R, Σ) is in ℓ -Update Inefficiency Normal Form (UINF) iff (R, Σ) is in ℓ -RFNF.

If (R, Σ) is in ℓ -RFNF (ℓ -UINF), then ℓ is a level of data redundancy (update inefficiency) that (R, Σ) prevents. Otherwise, (R, Σ) permits this level. If there is no ℓ for which (R, Σ) is in ℓ -RFNF (ℓ -UINF), no a priori upper bound exists for the level that (R, Σ) prevents. In this case, $\ell := \infty$.

4.3 ℓ -join efficiency

We define join efficiency as the maximum number of tuples that have matching values on the LHS of a non-trivial FD. This captures the update and join efficiency due to FDs, a core trade-off for logical schema design.

Formally, if r satisfies the non-trivial FD $X \rightarrow Y$ over relation schema R , then r is the lossless join of its projections on XY and $X(R - Y)$. That is, $r = \pi_{XY}(r) \bowtie \pi_{X(R-Y)}(r)$. Hence, the X -value

for a fixed tuple $t \in r$ joins the unique Y -value of t with any $R - Y$ -values of all tuples $t'_1, \dots, t'_\ell \in r$ that have matching values with t on X . The number ℓ of those tuples denotes the join strength of $t \in r$ for $X \rightarrow Y$. The join strength of r is the maximum join strength of any tuple in r . Finally, the join efficiency of (R, Σ) is the maximum join strength of any relation over R that satisfies Σ . A schema (R, Σ) is in ℓ -Join Efficiency Normal Form (JENF) iff the join efficiency of (R, Σ) is at most ℓ .

For example, the join efficiency of schema $(ECT, \{CT \rightarrow E, E \rightarrow C, \text{card}(E) \leq 1k\})$ is $1k$. The following relation r joins the redundant C -value *Kilo* with the $1k$ different values t_1, \dots, t_{1k} . As r satisfies $E \rightarrow C$, we have $r = r[EC] \bowtie r[ET]$.

$$\begin{array}{ccc} \begin{array}{ccc} E & C & T \\ \hline \text{Party} & \text{Kilo} & t_1 \\ \vdots & \vdots & \vdots \\ \text{Party} & \text{Kilo} & t_{1k} \end{array} & = & \begin{array}{cc} E & C \\ \hline \text{Party} & \text{Kilo} \end{array} \bowtie \begin{array}{cc} E & T \\ \hline \text{Party} & t_1 \\ \vdots & \vdots \\ \text{Party} & t_{1k} \end{array} \end{array}$$

If (R, Σ) is in ℓ -JENF, then ℓ is a level of join efficiency that (R, Σ) permits. Otherwise, ℓ is a level of join efficiency that (R, Σ) prevents. If there is no ℓ for which (R, Σ) is in ℓ -JENF, no a priori upper bound exists for the level of join efficiency that (R, Σ) permits. In this case, $\ell := \infty$.

4.4 Justification

It turns out for every $\ell \in \mathbb{N}_{\geq 1}^{\infty}$ that ℓ -BCNF coincides with ℓ -RFNF (ℓ -UINF) and with ℓ -JENF. Hence, schemata in ℓ -BCNF capture instances that are i) free from any ℓ -redundant data value occurrences (ℓ -update inefficiencies), and ii) permit level ℓ -join efficiency.

THEOREM 4.1. *For all (R, Σ) and all $\ell \in \mathbb{N}_{\geq 1}^{\infty}$, the following are equivalent: (1) (R, Σ) is in ℓ -RFNF (ℓ -UINF), (2) (R, Σ) is in ℓ -JENF, and (3) (R, Σ) is in ℓ -BCNF.*

The proof of Theorem 4.1 - in particular - constructs for every schema that is not in ℓ -BCNF an instance that features some ℓ -redundant data value occurrence (ℓ -update inefficiency), and a level $\ell + 1$ -join strength. Such a construction can be used in practice to automatically generate relations which exemplify the properties of a schema.

(HAP, Σ) from Example 3.2 is in ℓ_2 -BCNF but not in $\ell_2 - 1$ -BCNF. The CC $\text{card}(V) \leq \ell_2 - 1$ is not implied by Σ but $V \rightarrow C$ is an FD in Σ . Based on this FD we construct a relation with ℓ_2 tuples t'_1, \dots, t'_{ℓ_2} that all have matching values on columns in $V_{[\Sigma \text{FD}]}^+ = VC$, and unique values on all other columns. The relation may consist of the last ℓ_2 tuples in r of Table 1a. It prevents level ℓ_2 data redundancy and update inefficiency, and permits level ℓ_2 join efficiency.

5 SCHEMA DESIGN ALGORITHMS

After defining the update inefficiency and join efficiency of schema decompositions we propose three algorithms for logical schema design. We also illustrate how our concepts inform view selection already during logical design.

5.1 Assessing Decompositions

We can assess the quality of a schema decomposition \mathcal{D} for both updates and joins. The join efficiency of \mathcal{D} results from only those FDs preserved by \mathcal{D} . These may have been transformed into keys or not. We define the set of *join-supportive attribute subsets* as

$$JS_{\mathcal{D}}^{(R, \Sigma)} = \{S : X_k \mid \exists S \in \mathcal{D} \exists X \rightarrow Y \in \Sigma[S], \Sigma \models X \rightarrow S\} \cup \{S : X \mid \exists S \in \mathcal{D} \exists X \rightarrow Y \in \Sigma[S], \Sigma \not\models X \rightarrow S\}.$$

Next, the join strength of a join-supportive attribute set is given by the minimal upper bounds for any CC that applies to it. Hence, for (R, Σ) and $X \subseteq R$, $\ell_X := \min\{\ell \mid \Sigma \models \text{card}(X) \leq \ell\}$ is the minimum level of data redundancy for X . Now we define the level of join efficiency for \mathcal{D} as the maximum among all minimum levels of data redundancy for a join-supportive attribute set, that is,

$$\ell_{\mathcal{D}}^J := \sup\{\ell_X \mid S : X \in JS_{\mathcal{D}}^{(R, \Sigma)}\} \cup \{1 \mid S : X_k \in JS_{\mathcal{D}}^{(R, \Sigma)}\}.$$

In essence, FDs $X \rightarrow Y$ transformed into keys X contribute level 1, and other FDs contribute level ℓ_X . Defining $\ell_{\mathcal{D}}^J$ as supremum means it is 1 when Σ contains no FDs. In addition, the *total level of join efficiency for \mathcal{D}* is the sum of the levels:

$$\ell_{\mathcal{D}}^{J, \text{total}} = \sum_{S: X \in JS_{\mathcal{D}}^{(R, \Sigma)}} \ell_X + \sum_{S: X_k \in JS_{\mathcal{D}}^{(R, \Sigma)}} 1.$$

The update inefficiency of \mathcal{D} is determined by all FDs of the input. In particular, any lost FD will need to be enforced by joining elements of \mathcal{D} whenever updates occur. A BCNF decomposition of our running example into R_2 and R_3 results in a loss of the FD $E \rightarrow C$. As the instances r_2 and r_3 in Table 1b illustrate, any update of a C -value (such as *Kilo*) in r_2 needs to be propagated consistently for all its occurrences in r_2 with the same event (eg. *Party*), so the FD $E \rightarrow C$ still holds (on the join of R_2 and R_3). The example illustrates how lost FDs cause data redundancy on the join of schemata. This is the reason why they need to be enforced. Hence, we define the set of *update-critical attribute subsets* as

$$UC_{\mathcal{D}}^{(R, \Sigma)} = JS_{\mathcal{D}}^{(R, \Sigma)} \cup \{X \subseteq R \mid \exists X \rightarrow Y \in (\Sigma - (\cup_{S \in \mathcal{D}} \Sigma[S])^+)\}.$$

Now we define the level of update inefficiency for \mathcal{D} as the maximum among all minimum levels of data redundancy for an update-critical attribute set, that is,

$$\ell_{\mathcal{D}}^U := \sup\{\ell_X \mid S : X \in UC_{\mathcal{D}}^{(R, \Sigma)}\} \cup \{1 \mid S : X_k \in UC_{\mathcal{D}}^{(R, \Sigma)}\}.$$

Alternatively, we can also sum up to derive the *total level of update inefficiency for \mathcal{D}* as follows:

$$\ell_{\mathcal{D}}^{U, \text{total}} = \sum_{S: X \in UC_{\mathcal{D}}^{(R, \Sigma)}} \ell_X + \sum_{S: X_k \in UC_{\mathcal{D}}^{(R, \Sigma)}} 1.$$

With only FDs, BCNF decomposition or 3NF synthesis may lead to the optimal case of an FD-preserving BCNF decomposition with a join efficiency and update inefficiency of 1. Otherwise, some FD is lost or not a key. Hence, every non-optimal case results in unbounded levels of update inefficiency and join efficiency. With CCs, our framework can measure update inefficiency and join efficiency in all cases.

We will now propose different algorithms for schema design, and illustrate their achievements later by experiments.

Algorithm 2 OPT(R, Σ)

Require: (R, Σ) with CC/FD set Σ over schema R
Ensure: Lossless, FD-preserving $\ell_{\mathcal{D}}$ -BCNF decomposition \mathcal{D} of (R, Σ) with minimum $\ell_{\mathcal{D}}$ ($= \ell_{\mathcal{D}}^U = \ell_{\mathcal{D}}^J$)

- 1: Choose an atomic cover Σ_a of Σ [25];
- 2: **for all** $X \rightarrow A \in \Sigma_a$ **do**
- 3: $\Sigma[X \rightarrow A] \leftarrow \emptyset$
- 4: **for all** $Y \rightarrow B \in \Sigma_a$ ($YB \subseteq XA \wedge XA \not\subseteq Y_{[\Sigma]}^+$) **do**
- 5: $\ell_Y \leftarrow \min\{\ell \mid \Sigma \models \text{card}(Y) \leq \ell\}$
- 6: $\Sigma[X \rightarrow A] \leftarrow \Sigma[X \rightarrow A] \cup \{(Y \rightarrow B, \ell_Y)\}$
- 7: $\ell_{X \rightarrow A} \leftarrow \sup\{\ell_Y \mid Y \rightarrow B \in \Sigma[X \rightarrow A]\}$
- 8: $\mathcal{D} \leftarrow \emptyset$;
- 9: **for all** $X \rightarrow A \in \Sigma_a$ in decreasing order of $\ell_{X \rightarrow A}$ **do**
- 10: **if** $\Sigma_a - \{X \rightarrow A\} \models X \rightarrow A$ **then**
- 11: $\Sigma_a \leftarrow \Sigma_a - \{X \rightarrow A\}$
- 12: **else**
- 13: $\mathcal{D} \leftarrow \mathcal{D} \cup \{(XA, \Sigma_a[XA])\}$
- 14: Remove all $(S, \Sigma_a[S]) \in \mathcal{D}$ if $\exists(S', \Sigma_a[S']) \in \mathcal{D}$ ($S \subseteq S'$)
- 15: **if** there is no $(R', \Sigma') \in \mathcal{D}$ where $R' \rightarrow R \in \Sigma_{\mathcal{D}}^+$ **then**
- 16: Choose a minimal key K for R with respect to Σ
- 17: $\mathcal{D} \leftarrow \mathcal{D} \cup \{(K, \Sigma_a[K])\}$
- 18: **Return**($\mathcal{D}, \ell_{\mathcal{D}}$)

5.2 FD-preservation

Preserving all FDs guarantees that the levels of join efficiency and update inefficiency coincide. Otherwise, the level of update inefficiency may exceed that of join efficiency (more pain than gain). Consequently, we aim at minimizing the level of update inefficiency across all FD-preserving decompositions. Hence, we start with the unique *atomic cover* [25] (all L-reduced FDs with singleton attribute on the right-hand side), then compute for all FDs σ the maximum ℓ_{σ} that is associated with all the non-key FDs subsumed by σ , and then synthesize the final decomposition with schemata generated by FDs σ in decreasing order of their ℓ_{σ} , unless such a σ is implied by the remaining FDs. Due to Theorem 3.7, Algorithm 2 is worst-case exponential.

(HAP, Σ) from Example 3.2 is already in 3NF. As no FD-preserving BCNF decomposition exists, classical normalization stops here. Using (HAP, Σ) as input for Algorithm 2, we obtain $\Sigma_a = \Sigma \cup \{CT \rightarrow E\}$ (line 1). Hence, $\ell_{CT \rightarrow V} = 1m$ and $\ell_{CT \rightarrow E} = 1k$, and for other $X \rightarrow A \in \Sigma_a$, $\ell_{X \rightarrow A} = 1$ (line 7). However, $CT \rightarrow V$ is implied by $\Sigma_a - \{CT \rightarrow V\}$ (line 10). This step is critical, since $CT \rightarrow V$ has effectively been replaced by $CT \rightarrow E$, which also means that the 1m-level data redundancy caused by $CT \rightarrow V$ has been reduced to the 1k-level data redundancy caused by $CT \rightarrow E$. Furthermore, $R_1 \subseteq R_4$ (line 13), so Algorithm 2 returns the 1k-BCNF schema $\mathcal{D}_1 = \{(R_2, \Sigma_2), (R_3, \Sigma_3), (R_4, \Sigma_4)\}$ from the intro. Classical normalization would generate neither \mathcal{D}_1 nor \mathcal{D}_2 . Even if it did, \mathcal{D}_1 , \mathcal{D}_2 and (HAP, Σ) would be assessed as equal in quality.

We may apply Algorithm 2 to a partial input of Σ . For example, excluding FDs which potentially cause high update inefficiency but are unlikely violated by actual updates, may result in higher join efficiency. Similarly, one may include all FDs $X \rightarrow Y$ where ℓ_X meets a given threshold.

Algorithm 3 GREED(R, Σ)

Require: (R, Σ) with CC/FD set Σ over R , $\ell \in \mathbb{N}_{\geq 1}$
Ensure: Lossless $\ell_{\mathcal{D}}^U$ -BCNF decomposition \mathcal{D} of (R, Σ)

- 1: **if** (R, Σ) is in ℓ -BCNF **then**
- 2: $\mathcal{D} \leftarrow \{(R, \Sigma)\}$
- 3: **else**
- 4: $\ell_{\max} \leftarrow \max\{\ell_X \mid X \rightarrow Y \in \Sigma, Y \not\subseteq X, \Sigma \not\models X \rightarrow R\}$
- 5: Pick $X \rightarrow Y \in \Sigma$ ($Y \not\subseteq X \wedge \Sigma \not\models X \rightarrow R \wedge \ell_X = \ell_{\max}$)
- 6: $R_1 \leftarrow XY; R_2 \leftarrow X(R - XY)$
- 7: $\mathcal{D} \leftarrow \text{GREED}(R_1, \Sigma[R_1]) \cup \text{GREED}(R_2, \Sigma[R_2])$
- 8: **Return**($\mathcal{D}, \ell_{\mathcal{D}}^U, \ell_{\mathcal{D}}^J$)

5.3 Greedy Decomposition

Typically, FD-preservation requires many output tables, which is linked to low join efficiency. We reposition classical BCNF decomposition as a greedy algorithm for reducing $\ell_{\mathcal{D}}^U$ using fewer tables. Selecting non-key FDs that drive classical BCNF decomposition is arbitrary. With CCs at hand, we select FDs whose level of data redundancy is maximal (for example, ∞). Algorithm 3 shows this strategy. Targeting higher join efficiency, we may use a higher level ℓ as input. If unspecified, the default target is $\ell = 1$. The algorithm proceeds with some best remaining FD (lines 4-5), as long as some local schema is not in ℓ -BCNF (line 1). The output returns schema \mathcal{D} , the level $\ell_{\mathcal{D}}^U$ of update inefficiency it prevents, and the level $\ell_{\mathcal{D}}^J$ of join efficiency it permits. The penalty for fewer output tables is the lack of control over these levels due to lost FDs. FDs $X \rightarrow A$ where $\ell_X = \infty$ have high priority. With no CCs given, we have the classical BCNF decomposition case where $\ell_{\mathcal{D}}^U = \infty = \ell_{\mathcal{D}}^J$, unless \mathcal{D} is FD-preserving and $\ell_{\mathcal{D}}^U = 1 = \ell_{\mathcal{D}}^J$.

We apply Algorithm 3 to (HAP, Σ) from Example 3.2. As $V \rightarrow C \in \Sigma$ causes $\ell_{\max} = 1m$ (line 4), we obtain $\mathcal{D}_g = \{(R_2 = VC, \Sigma_2 = \{V \rightarrow C\}), (R_3 = EVT, \Sigma_3 = \{TV \rightarrow E, ET \rightarrow V\})\}$ (line 7). We lost $E \rightarrow C$ and $CT \rightarrow V$, so $\ell_{\mathcal{D}_g}^U = 1k$. We create assertion checks to enforce the lost FDs $E \rightarrow C$ and $CT \rightarrow V$ on the join $R_2 \bowtie R_3$:

```
CREATE ASSERTION LostFDEtoC CHECK( NOT EXISTS (
  SELECT R3.E FROM R2, R3 WHERE R2.V = R3.V
  GROUP BY R3.E HAVING COUNT(R2.C)>1);
CREATE ASSERTION LostFDCTtoV CHECK( NOT EXISTS (
  SELECT R2.C, R3.T FROM R2, R3 WHERE R2.V = R3.V
  GROUP BY R2.C, R3.T HAVING COUNT(R2.V)>1);
```

\mathcal{D}_g achieves the same level of update inefficiency as \mathcal{D}_1 with fewer schemata. This comes at the price of lost FDs that require assertion checks and smaller join efficiency $\ell_{\mathcal{D}_g}^J = 1$.

Our algorithms can be adapted when more information becomes available. For example, while \mathcal{D}_g was driven by the FD $V \rightarrow C$ due to $\text{card}(V) \leq 1m$, there may not be updates of C -values based on V -values but frequent updates of C -values based on E -values. Hence, we could prioritize the FD $E \rightarrow C$ instead to obtain $\mathcal{D}'_g = \{(R_1 = EC, \Sigma_1 = \{E \rightarrow C\}), (R_3, \Sigma_3)\}$ with lost FDs $V \rightarrow C$ and $CT \rightarrow V$.

Just like FD-preserving BCNF decompositions cannot always be achieved, ℓ -BCNF cannot always be achieved for a given ℓ . Adapting a classical example [3], consider $R = \{c(\text{ity}), s(\text{street}), z(\text{ip}), A\}$ where

Algorithm 4 HYBRID(R, Σ)

Require: (R, Σ) with CC/FD set Σ over schema R
Ensure: Lossless $\ell_{\mathcal{D}}^U$ -BCNF decomposition \mathcal{D} of (R, Σ)

- 1: $\mathcal{D} \leftarrow \text{OPT}(R, \Sigma)$
- 2: **for all** $S \in \mathcal{D}$ **do**
- 3: $S = XA$ for some $X \rightarrow A \in \Sigma_a$
- 4: **for all** $Y \rightarrow B \in \Sigma_a(YB \subseteq XA \wedge XA \not\subseteq Y_{\Sigma[FD]}^+)$ **do**
- 5: **if** $\ell_Y > \ell_X$ **then**
- 6: $S_1 = YB; S_2 = Y(XA - B)$
- 7: $\mathcal{D} \leftarrow (\mathcal{D} - \{(S, \Sigma_a[S])\}) \cup \{(S_1, \Sigma_a[S_1]), (S_2, \Sigma_a[S_2])\}$
- 8: Remove any non-maximal schema from \mathcal{D}
- 9: **Return**($\mathcal{D}, \ell_{\mathcal{D}}^U, \ell_{\mathcal{D}}^J$)

Σ consist of $sc \rightarrow z$ and $z \rightarrow c$, and the CCs $\text{card}(c, s) \leq \ell + 1$ and $\text{card}(z) \leq \ell + 1$. If $\{c, s, z\}$ is not in the output \mathcal{D} , we lose $sc \rightarrow z$, and $\ell_{\mathcal{D}}^U = \ell + 1$. If $\{c, s, z\}$ is in \mathcal{D} , we will have $z \rightarrow c$ on $\{c, s, z\}$, and hence $\ell_{\mathcal{D}}^U = \ell + 1$. Thus, no output can be in ℓ -BCNF.

5.4 Hybrid algorithm

We combine the two previous strategies to the hybrid Algorithm 4. Here, Algorithm 2 is applied first (line 1) to obtain the smallest possible $\ell_{\mathcal{D}}$ among all FD-preserving decompositions. Next we check if for any of the output schemata XA (lines 2-3) that are in 3NF but not in BCNF (line 4), any non-key FD $Y \rightarrow B$ on XA satisfies $\ell_Y > \ell_X$ (line 5). In this case, we trade the preservation of $X \rightarrow A$ for the lower level ℓ_X of update inefficiency by decomposing with the FD $Y \rightarrow B$ (lines 6/7).

On input (HAP, Σ) from Example 3.2, Algorithm 4 returns decomposition \mathcal{D}_1 as output of Algorithm 2 (line 1). In particular, $S = R_4 = CET$ is in 3NF but not in BCNF for $\Sigma_a[S] = \Sigma_4 = \{CT \rightarrow E, E \rightarrow C\}$. That is, the FD $E \rightarrow C \in \Sigma_a$ is critical on the schema $S = R_4 = CET$ (lines 3 and 4). Since $\ell_E = 1k > 1 = \ell_{CT}$ (line 5), we decompose along the critical FD $E \rightarrow C$ to replace R_4 by $R_6 = EC$ and $R_7 = ET$ (lines 6 and 7). However, $R_6 = R_1$ and $R_7 \subseteq R_3$. Hence, the output is $\mathcal{D}_h = \{(R_1, \Sigma_1), (R_2, \Sigma_2), (R_3, \Sigma_3)\}$ with $\ell_{\mathcal{D}_h}^U = 1 = \ell_{\mathcal{D}_h}^J$ and lost FD $CT \rightarrow V$. The assertion check LostFDCTtoV enforces the FD $CT \rightarrow V$ on the join $R_2 \bowtie R_3$.

Table 4 summarizes the schemata that our algorithms return for our example, and also \mathcal{D}_2 , including their properties.

5.5 Materialized Views

When operational, frequent access patterns emerge on the database. Adding materialized views can help accelerate queries but will occupy additional storage space and time to maintain data consistency. Hence, view selection should be informed by quantifying join support and maintenance costs. We define the (total) level of join efficiency and update inefficiency of a given view \mathcal{V} (namely, a set of attributes) as $\ell_{\mathcal{V}}^{J, \text{total}}$ and $\ell_{\mathcal{V}}^{U, \text{total}}$, respectively.

For our hybrid schema \mathcal{D}_h as example, a frequent query may ask at what times companies work on a venue. Hence, we may introduce the following materialized view \mathcal{V} :

```
CREATE MATERIALIZED VIEW  $\mathcal{V}$  AS (
  SELECT R2.V, R2.C, R3.T FROM R2, R3 WHERE R2.V = R3.V);
```

Table 4: Schema Designs for Running Example

Method	Schema	Lost FDs	ℓ^U	ℓ^J	Materialized View
OPT	$\mathcal{D}_1: R_2 = CV \text{ with } V \rightarrow C$ $R_3 = ETV \text{ with } TV \rightarrow E, ET \rightarrow V$ $R_4 = CET \text{ with } CT \rightarrow E, E \rightarrow C$	none	1k	1k	$\mathcal{V}_{\mathcal{D}_1} = \pi_{VCT}(R_2 \bowtie R_3)$
N/A	$\mathcal{D}_2: R_1 = EC \text{ with } E \rightarrow C$ $R_3 = ETV \text{ with } TV \rightarrow E, ET \rightarrow V$ $R_5 = CTV \text{ with } CT \rightarrow V, V \rightarrow C$	none	1m	1m	$\mathcal{V}_{\mathcal{D}_2} = \pi_{ECT}(R_1 \bowtie R_3)$
GREED	$\mathcal{D}_g: R_2 = CV \text{ with } V \rightarrow C$ $R_3 = ETV \text{ with } TV \rightarrow E, ET \rightarrow V$	$E \rightarrow C$ $CT \rightarrow V$	1k	1	$\mathcal{V}_{\mathcal{D}_g} = \pi_{ECT}(R_2 \bowtie R_3)$
HYBRID	$\mathcal{D}_h: R_1 = EC \text{ with } E \rightarrow C$ $R_2 = CV \text{ with } V \rightarrow C$ $R_3 = ETV \text{ with } TV \rightarrow E, ET \rightarrow V$	$CT \rightarrow V$	1	1	$\mathcal{V}_{\mathcal{D}_h} = \pi_{VCT}(R_1 \bowtie R_3)$

Table 5: Queries q_1 and q_2 on Schema Designs for Running Example

q	\mathcal{D}_1	$\mathcal{D}_1 + \mathcal{V}_{\mathcal{D}_1}$	\mathcal{D}_2	$\mathcal{D}_2 + \mathcal{V}_{\mathcal{D}_2}$	\mathcal{D}_g	$\mathcal{D}_g + \mathcal{V}_{\mathcal{D}_g}$	\mathcal{D}_h	$\mathcal{D}_h + \mathcal{V}_{\mathcal{D}_h}$
q_1	$\pi_{ECT}(R_4)$	$\pi_{ECT}(R_4)$	$\pi_{ECT}(R_1 \bowtie R_3)$	$\mathcal{V}_{\mathcal{D}_2}$	$\pi_{ECT}(R_2 \bowtie R_3)$	$\mathcal{V}_{\mathcal{D}_g}$	$\pi_{ECT}(R_1 \bowtie R_3)$	$\pi_{ECT}(R_1 \bowtie R_3)$
q_2	$\pi_{CTV}(R_2 \bowtie R_3)$	$\mathcal{V}_{\mathcal{D}_1}$	$\pi_{CTV}(R_5)$	$\pi_{CTV}(R_5)$	$\pi_{CTV}(R_2 \bowtie R_3)$	$\pi_{CTV}(R_2 \bowtie R_3)$	$\pi_{CTV}(R_1 \bowtie R_3)$	$\mathcal{V}_{\mathcal{D}_h}$

In this case, the assertion check LostFDCTtoV can directly be enforced on the view instead.

```
CREATE ASSERTION LostFDCTtoV CHECK(NOT EXISTS
  (SELECT C, T FROM  $\mathcal{V}$ 
   GROUP BY C, T HAVING COUNT(V)>1));
```

Due to $V \rightarrow C, CT \rightarrow V$ and $\text{card}(V) \leq 1m$, the view has level $\ell_V^J = 1m = \ell_V^U$ update inefficiency and join efficiency, respectively. Join support is effective, but requires the propagation of a C -update on base table R_2 to up to $1m$ updates on \mathcal{V} . Hence, the trade-off we quantify between update inefficiency and join efficiency is also intrinsic to materialized views and should be part of the information for helping with their selection. Table 4 also lists some views for the various schema designs of our running example.

6 EXPERIMENTS

Our experiments analyze the properties of our normal forms and how these translate into the performance of updates and joins over instances of the schema. We implemented our algorithms in Visual C++. Experiments were done on an Intel Xeon W-2123, 3.6 GHz, 256GB, Windows 10 PC, with the 2017 SQL Server Community Edition.

6.1 Qualitative study

Firstly, we analyze the performance of two updates and two joins on synthetic instances over the schema designs for our running example.

Our first instance over (HAP, Σ) consists of 1,001,000 tuples. There are $1k$ tuples with matching values on E and C , and $1m$ tuples with matching values on V and C . The other values are unique in their columns. The instance is isomorphic to instance r in Table 1a with $\ell_1 = 1k$ and $\ell_2 = 1m$. To illustrate the impact of FDs on updates and joins, the numbers of redundant values they cause in the instance coincides with their levels of data redundancy on the schema (ℓ_1 and ℓ_2).

For the experiments, we consider four operations that are affected by our CCs and FDs:

- Update u_1 replaces all occurrences of a C -value associated with a given E -value,
- Update u_2 replaces all occurrences of a C -value associated with a given V -value,
- Query q_1 returns all CTE combinations, and
- Query q_2 returns all CTV combinations.

Each operation is run 100 times on each of the schema designs in Table 4, and the average value reported.

Updates are always propagated to views from base tables. As $E \rightarrow C$ is lost on \mathcal{D}_g , u_1 updates C -values on $R_2 \bowtie R_3$ and projects them onto R_2 . Table 5 shows how the queries q_1 and q_2 look like on each of our schemata from Table 4, without and with materialized views.

On the instance over 3NF schema (HAP, Σ) , q_1 and q_2 are simple projections and do not require joins. For u_1 , $1k$ occurrences of a redundant C -value are updated, and $1m$ occurrences of a redundant C -value are updated for u_2 , due to the non-key FDs $E \rightarrow C$ and $V \rightarrow C$, respectively.

Table 6 shows the times (in seconds) for the operations on the instance projected onto the schemata of Table 4.

Specific updates and queries are best supported by different designs. Notably, the level of data redundancy for an FD translates proportionally into the performance for updates and joins affected by the FD. This also applies to the views.

Operations u_2 and q_1 are fast on \mathcal{D}_1 with R_4 and key V on R_2 , but slow on \mathcal{D}_2 . Symmetrically, u_1 and q_2 are fast on \mathcal{D}_2 with R_5 and key E on R_1 , but slow on \mathcal{D}_1 . For each operation, the difference in performance between \mathcal{D}_1 and \mathcal{D}_2 is proportional to the level of data redundancy caused by the FD that affects the operation. For example, the FD $V \rightarrow C$ with $\ell_V = 1m$ causes a significant difference between performing u_2 (44.58s difference) on \mathcal{D}_1 and \mathcal{D}_2 , and between performing q_2 (1.35s difference) on \mathcal{D}_1 and \mathcal{D}_2 . In

Table 6: Synthetic Experiment 1 (in seconds)

Op	HAP	\mathcal{D}_1	$\mathcal{D}_1 + \mathcal{V}_{\mathcal{D}_1}$	\mathcal{D}_2	$\mathcal{D}_2 + \mathcal{V}_{\mathcal{D}_2}$	\mathcal{D}_g	$\mathcal{D}_g + \mathcal{V}_{\mathcal{D}_g}$	\mathcal{D}_h	$\mathcal{D}_h + \mathcal{V}_{\mathcal{D}_h}$
u_1	0.93	0.91	0.91	0.59	1.03	1.62	1.81	0.59	0.59
u_2	45.75	0.23	45.33	44.81	44.81	0.20	0.20	0.21	83.82
q_1	0.011	0.009	0.009	0.17	0.012	1.32	0.009	0.14	0.14
q_2	3.39	4.11	3.47	2.76	2.76	4.09	4.09	4.32	3.44

Table 7: Synthetic Experiment 2 (in milliseconds)

Op	HAP	\mathcal{D}_1	$\mathcal{D}_1 + \mathcal{V}_{\mathcal{D}_1}$	\mathcal{D}_2	$\mathcal{D}_2 + \mathcal{V}_{\mathcal{D}_2}$	\mathcal{D}_g	$\mathcal{D}_g + \mathcal{V}_{\mathcal{D}_g}$	\mathcal{D}_h	$\mathcal{D}_h + \mathcal{V}_{\mathcal{D}_h}$
u_1	58.2	55.8	55.8	3.7	60.7	82.3	170.2	3.7	3.7
u_2	10	4.4	12.7	9.8	9.8	4.9	4.9	4.7	18.9
q_1	8.3	7.5	7.5	11.1	10.6	236	4	11.6	11.6
q_2	5.6	14.6	5.7	5.3	5.3	14.5	14.5	14.1	7.4

comparison, the FD $E \rightarrow C$ with $\ell_V = 1k$ causes a much smaller difference between performing u_1 (0.32s) on \mathcal{D}_1 and \mathcal{D}_2 , and between performing q_1 (0.161s) on \mathcal{D}_1 and \mathcal{D}_2 .

Relatively to $\ell_{\mathcal{V}_{\mathcal{D}_1}} = 1m$, the speed up of q_2 and slow down of u_2 by $\mathcal{V}_{\mathcal{D}_1}$ are large. Relatively to $\ell_{\mathcal{V}_{\mathcal{D}_2}} = 1k$, the speed up of q_1 and slow down of u_1 by $\mathcal{V}_{\mathcal{D}_2}$ are smaller.

As expected, \mathcal{D}_g performs well on u_2 but not so much on the other operations. $\mathcal{V}_{\mathcal{D}_g}$ speeds up q_1 with small penalty for u_1 , but which had poor performance already.

\mathcal{D}_h performs well for both u_1 and u_2 , and quite well for q_1 . $\mathcal{V}_{\mathcal{D}_h}$ speeds up q_2 but slows down u_2 significantly.

Table 7 shows our results for the second synthetic instance with 1,100 tuples, $\ell_1 = 1k$, and $\ell'_2 = 100 < 1m$. The experiment illustrates that the same combinations of operations are best supported by the same designs, when compared to the first instance. Since the actual cardinality (100) is much smaller than what is permitted ($1m$), processing times for operations with this cardinality are much smaller (in ms).

The choice of a logical design depends on how well it is believed to support the future workload of operations. Our framework quantifies this support by the level of data redundancy that FDs exhibit as part of the design or materialized view, which measures their impact on the performance of specific updates and joins. In contrast, classical logical design may not bring forward some designs and cannot quantify the support for workloads. For example, schema (HAP, Θ) from Example 3.2 is in 3NF and no lossless, FD-preserving decomposition into BCNF exists. Hence, classical design based on FDs alone would stop there. In contrast, by including CCs in the analysis, our methods transform the schema (HAP, Σ) from Example 3.2 into the various schema designs listed in Table 4 and quantify the achievements in terms of update inefficiency and join efficiency.

6.2 Quantitative study

Secondly, we analyze update and join performance over schema designs derived by our algorithms for FDs and CCs we mined from the real-world data sets in Table 8. It shows the number of FDs discovered (#FD), the total number of null marker occurrences (# \perp), and the number of rows and columns (#R, #C).

Table 8: Data Sets Used in Experiments

Data set	#FDs	# \perp	#R	#C
china	918	418580	262920	18
ncvoter	568	3079822	1024000	19

The data sets benchmark algorithms that discover FDs $X \rightarrow A$ from data and rank them in decreasing order by the lowest bound ℓ_X such that $card(X) \leq \ell_X$ holds on the data [41, 42, 48]. We regard different occurrences of \perp in the same column as matching domain values ($\perp = \perp$). Our results are very similar otherwise ($\perp \neq \perp$). The discovered CCs and FDs represent patterns on the data, but not necessarily on the domain. Input to our algorithms are only more relevant FDs, namely those ranked in the top-20%. Adding further FDs causes hardly any differences, indicating that FDs with higher levels of data redundancy determine schema designs.

6.2.1 Analysis of update and join efficiency levels. We applied Algorithms 2 (OPT), 3 (GREED), and 4 (HYBRID) to the schemata for *ncvoter* and *china*.

Table 9 shows the levels $\ell^U; \ell^J$ for the decompositions \mathcal{D} and the original input, their total levels $\ell^{U,\text{total}}; \ell^{J,\text{total}}$, output size $|\mathcal{D}|$, and the runtime.

The reduction in update inefficiency levels is significant for OPT and HYBRID. GREED achieves small reductions due to lost FDs. The join efficiency level for GREED equals 1 since every FD becomes either a key or is lost. Under FD-preservation, OPT always achieves the optimum level ℓ^U , which always equals ℓ^J as all FDs are preserved. HYBRID achieves even lower levels ℓ^U by loss of some FDs and lowering ℓ^J .

The reductions of ℓ^U are achieved by different output sizes: GREED requires few schemata, while OPT achieves the optimum by adding schemata to preserve all FDs, and HYBRID achieves further reductions by even more schemata. All algorithms are fast as the input is an atomic cover [25].

6.2.2 Physical Performance. We ran updates and join queries over our original data sets and the outputs of Algorithms 2 and 3 on input sets that included all FDs $X \rightarrow Y$ where ℓ_X met a given target level ℓ . Figures 3 and 5 show the results of Algorithm 2, and

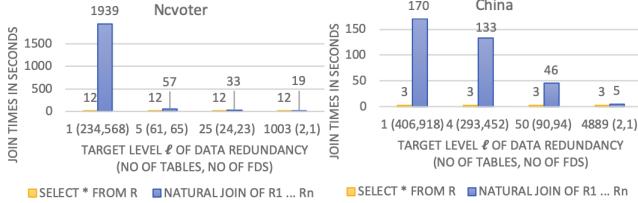
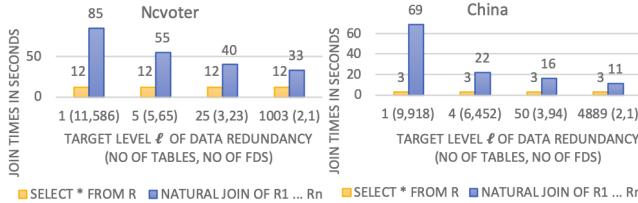
Table 9: Algorithmic Achievements on Real Schemata

measures	ORIGINAL	OPT	GREED	HYBRID
ℓ^U	4,097	15	512	7
ℓ^J	4,097	15	1	7
$\ell^{U,\text{total}}$	5,043	87	3,035	79
$\ell^{J,\text{total}}$	5,043	87	6	73
$ \mathcal{D} $	1	65	6	66
time (ms)	-	< 1	< 1	< 1

(a) Measures for top-20% FDs of ncvoter

measures	ORIGINAL	OPT	GREED	HYBRID
ℓ^U	4,889	56	4,881	55
ℓ^J	4,889	56	1	50
$\ell^{U,\text{total}}$	62,492	327	54,247	257
$\ell^{J,\text{total}}$	62,492	327	4	165
$ \mathcal{D} $	1	145	4	157
time (ms)	-	< 1	< 1	< 1

(b) Measures for top-20% FDs of china

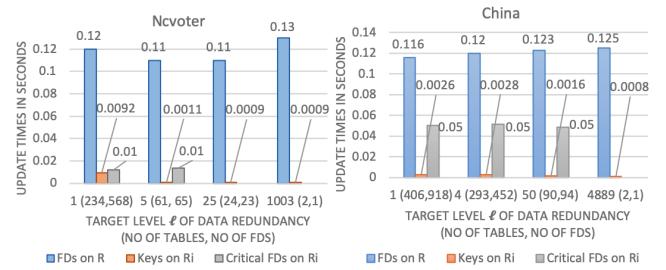
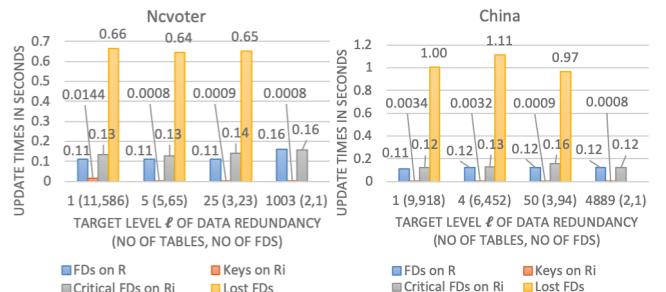
**Figure 3: Join performance after using OPT****Figure 4: Join performance after using GREED**

Figures 4 and 6 those of Algorithm 3. The x -axes always show the target levels ℓ , the number n of output tables and the number of input FDs. The y -axes in Figures 3 and 4 show the runtimes (in seconds) for a join query, and the runtimes (in seconds) for updates in Figures 5 and 6.

The join queries are:

- $\text{SELECT } * \text{ FROM HAP}$ over the original schema HAP, and
- $\text{SELECT } * \text{ FROM } R_1 \text{ NATURAL JOIN } \dots \text{ NATURAL JOIN } R_n \text{ over } \{R_1, \dots, R_n\}$.

Join times increase with lower target levels ℓ . For OPT, join performance becomes poor due to the many output tables required for FD preservation. This is compensated by GREED, but the trade-off is poor update performance for lost FDs. The level ℓ , solely determined by the input integrity constraints, translates into different physical performance degrees for updates and joins on the output

**Figure 5: Update performance after using OPT****Figure 6: Update performance after using GREED**

of our algorithms. Making these options available is the point of our new framework.

Updates involve changes to all occurrences of some redundant value (recall the update scenario from Section 4). After an FD is transformed into a key, every value occurs at most once, so an update is required for at most one occurrence.

Average update times for keys are about two orders of magnitude lower than those for input FDs on the original data. For preserved, non-key FDs (called critical FDs) the gain is about one order of magnitude (grey bars in Figures 5 and 6, with no data point when all FDs become keys). Lost FDs require more update time than on the original data set, due to necessary joins. With no lost FDs, only keys need to be enforced (in Figure 6 for the largest target level). Figures 5 and 6 show that, by lowering the level ℓ , more FDs are converted into keys (for OPT), which improves update efficiency at the expense of introducing more tables. A main message is that schema design is primarily driven by high-ranked FDs, but inclusion of low-ranked FDs means that more FDs need to be preserved during synthesis to lower their update time (by transformation into keys or critical FDs with OPT), or more FDs will be lost during decomposition (GREED).

6.2.3 Comments. Logical schema design is driven by business rules which constrain instances to those that are meaningful to the underlying domain. We have seen how the upper bounds of CCs inform schema design and quantify the support for query and update operations. Our experiments illustrate that smaller bounds mean smaller differences in performance between designs. During logical design the choice of a schema is thus determined by the upper bounds, as these represent the current requirements. When requirements change over time, physical database tuning, such as de-normalization, can lift performance. However, this is only

possible once the database is operational, when actual instances can be observed and stable patterns of workloads have emerged. If constraints evolve to the point that a different design is more suitable, then this may justify migration to a new schema. Ultimately, logical design informs the choice of a schema based on application requirements, independently of whether they have evolved or not.

Ranking designs informs the search for a relevant schema, just as ranking websites informs the search for relevant websites. Blindly picking a design with minimum ℓ is similar to using the “I’m feeling lucky” button. This also extends to the selection of materialized views. Our experiments suggest to carefully examine which operations are affected by which FDs. The level of data redundancy caused by the FDs quantifies how much they affect the operations. In particular, if the level of data redundancy caused by some FD is a priori unbounded ($\ell = \infty$), our framework alerts analysts to the fact that requirements analysis may still be incomplete.

7 RELATED WORK

Schema design has been studied in data models such as SQL [26], nested relations [39], data warehouses [30], object-orientation [24], semantics [9], temporality [23], the Web [1], uncertainty [35], and graphs [16]. Similar to classical design, our results may be extended to richer data models and higher normal forms such as 4NF [14], Project-Join Normal Form [15], or Inclusion-Dependency Normal Form [33].

Unlike BCNF [4, 5] and 3NF [6] which only use FDs, CCs measure the effort for achieving data consistency. Our algorithms return designs that quantify worst-case update inefficiency and best-case join efficiency. Our concepts also quantify incremental maintenance and join support of materialized views during logical design. In contrast, numerical dependencies (NDs) apply classical normalization after horizontal decomposition into blocks where FDs hold [19]. NDs are not finitely axiomatizable [18].

Schema design is not expected to optimize the layout for all instances. Instead, physical tuning kicks in after deploying the logical schema and once reliable data retrieval patterns emerge. These include virtual de-normalization in main memory databases [37], migration to NoSQL [49], and data warehouse designs [30]. Guidelines have been devised to recover 3NF from de-normalised schema [43], and de-normalizing normalised schemata [7]. Information-theoretic justifications exist for fact tables in snowflake schemata [32] and for 3NF [28]. Update inefficiency and join efficiency inform the difficult problem of selecting materialized views [11].

Kojić and Milićev de-normalize by maximizing read benefits while write costs remain under a threshold [27]. Their technique tunes a logical schema (for example, in 3NF) based on specific updates and queries. Hence, their approach is applicable once the database is operational and a mature workload model available. This additional input makes it possible to derive a schema optimized for the input workload for [27]. The work in [27] does not use CCs. Our approach is for logical schema normalization where a workload model in terms of frequent update and query operations is not available yet. We use CCs to explore various normal forms with different properties in terms of update inefficiency and join efficiency on the schemata. Hence, the design team can assess different designs based on their properties. We only require a set

of attributes, FDs, and CCs as input. As our work brings forward a logical schema, it may provide a different starting point for applying the work in [27] than classical normalization does. For example, once frequent updates and queries have emerged, we may apply [27] to the output of Algorithm 2 rather than the 3NF schema (HAP, Σ).

Recent approaches to schema design and evolution for NoSQL databases [45] are driven by a query workload, and NoSQL schemata are thus de-normalized. We are unaware of work for logical NoSQL schema design for CCs and FDs.

“Cardinality constraints are one of the most important kinds of constraint in conceptual modeling” [40]. “CCs correspond to very common semantic rules on relationships and their formal definition at the conceptual level improves significantly the completeness of data description” [31]. Surprisingly, our way of applying CCs to logical schema design has not been observed before. CCs were introduced in Chen’s seminal ER paper [9], and have been studied in data models such as semantic [34], Web [17], spatial and temporal [13], and uncertain models [44]. They are part of major languages for data and knowledge modeling, including UML, EER, ORM, XSD (such as `maxOccurs`), or OWL (such as `owl:maxCardinality`) [20]. They have also been used in data cleaning [10], database testing [8], query answering [12] and reverse engineering [46].

Violations of dependencies by dirty data motivate relaxed notions such as approximate keys and FDs [22, 29]. When mining from (dirty) data, approximate notions may improve the recall but worsen the precision of identifying meaningful rules. The CC $card(X) \leq \ell$ may be seen as an approximate key permitting up to ℓ duplicates. For small ℓ , we may therefore view our work as extending classical schema design to approximate keys, offering some robustness for dirty data.

8 CONCLUSION AND FUTURE WORK

We introduced the first logical schema design framework that measures update inefficiency and join efficiency, based on integrity constraints alone. This is possible by the new notion of level- ℓ data redundancy, which is determined by the upper bounds of CCs at schema design time. Our infinite family of ℓ -Bounded Cardinality Normal Forms characterizes instances that are free from level ℓ data redundancy and update inefficiency, and permit level ℓ join efficiency. We developed algorithms for schema design, and illustrated experimentally how they reduce the levels of update inefficiency and join efficiency with trade-offs in the size of output designs. We also showed experimentally how these levels quantify the suitability of schema designs and materialized views for the performance of specific updates and joins on instances of the designs. Our framework uses domain knowledge about CCs to advance logical schema design.

Future work will address more constraints and data models. We expect the interaction of CCs and join dependencies to challenge the development of higher normal forms.

ACKNOWLEDGMENTS

The authors thank the anonymous reviewers for their valuable suggestions. The first author thanks Joachim Biskup, Bernhard Thalheim, and Jef Wijsen for insightful discussions during Dagstuhl Seminar 19031.

REFERENCES

- [1] Marcelo Arenas. 2006. Normalization theory for XML. *SIGMOD Record* 35, 4 (2006), 57–64.
- [2] William Ward Armstrong. 1974. Dependency Structures of Data Base Relationships. In *IFIP congress*, Vol. 74. Geneva, Switzerland, 580–583.
- [3] Catriel Beeri and Philip A Bernstein. 1979. Computational problems related to the design of normal form relational schemas. *ACM Trans. Database Syst.* 4, 1 (1979), 30–59.
- [4] Catriel Beeri, Philip A Bernstein, and Nathan Goodman. 1978. A sophisticate's introduction to database normalization theory. In *VLDB*. 113–124.
- [5] Philip A. Bernstein and Nathan Goodman. 1980. What does Boyce-Codd Normal Form Do?. In *VLDB*. 245–259.
- [6] Joachim Biskup, Umeshwar Dayal, and Philip A. Bernstein. 1979. Synthesizing Independent Database Schemas. In *SIGMOD*. 143–151.
- [7] Douglas B. Bock and John F. Schrage. 2002. Denormalization guidelines for base and transaction tables. *ACM SIGKDD Bull.* 34, 4 (2002), 129–133.
- [8] Nicolas Bruno, Surajit Chaudhuri, and Dilys Thomas. 2006. Generating Queries with Cardinality Constraints for DBMS Testing. *IEEE Trans. Knowl. Data Eng.* 18, 12 (2006), 1721–1725.
- [9] Peter Chen. 1976. The entity-relationship model-toward a unified view of data. *ACM Trans. Database Syst.* 1, 1 (1976), 9–36.
- [10] Wenguang Chen, Wenfei Fan, and Shuai Ma. 2009. Incorporating cardinality constraints and synonym rules into conditional functional dependencies. *Inf. Process. Lett.* 109, 14 (2009), 783–789.
- [11] Rada Chirkova and Jun Yang. 2012. Materialized Views. *Found. Trends Databases* 4, 4 (2012), 295–405.
- [12] Graham Cormode, Divesh Srivastava, Entong Shen, and Ting Yu. 2012. Aggregate Query Answering on Possibilistic Data with Cardinality Constraints. In *ICDE*. 258–269.
- [13] Faiz Currim and Sudha Ram. 2008. Conceptually modeling windows and bounds for space and time in database constraints. *Commun. ACM* 51, 11 (2008), 125–129.
- [14] Ronald Fagin. 1977. Multivalued Dependencies and a New Normal Form for Relational Databases. *ACM Trans. Database Syst.* 2, 3 (1977), 262–278.
- [15] Ronald Fagin. 1979. Normal Forms and Relational Database Operators. In *SIGMOD*. 153–160.
- [16] Wenfei Fan, Yinghui Wu, and Jingbo Xu. 2016. Functional Dependencies for Graphs. In *SIGMOD*. 1843–1857.
- [17] Flavio Ferrarotti, Sven Hartmann, and Sebastian Link. 2013. Efficiency frontiers of XML cardinality constraints. *Data Knowl. Eng.* 87 (2013), 297–319.
- [18] John Grant and Jack Minker. 1985. Inferences for Numerical Dependencies. *Theor. Comput. Sci.* 41 (1985), 271–287.
- [19] John Grant and Jack Minker. 1985. Normalization and Axiomatization for Numerical Dependencies. *Inf. Control.* 65, 1 (1985), 1–17.
- [20] Neil Hall, Henning Köhler, Sebastian Link, Henri Prade, and Xiaofang Zhou. 2015. Cardinality constraints on qualitatively uncertain data. *Data Knowl. Eng.* 99 (2015), 126–150.
- [21] Sven Hartmann. 2003. Reasoning about participation constraints and Chen's constraints. In *ADC*. 105–113.
- [22] Ykä Huhtala, Juha Kärkkäinen, Pasi Porkka, and Hannu Toivonen. 1999. TANE: An Efficient Algorithm for Discovering Functional and Approximate Dependencies. *Comput. J.* 42, 2 (1999), 100–111.
- [23] Christian S. Jensen, Richard T. Snodgrass, and Michael D. Soo. 1996. Extending Existing Dependency Theory to Temporal Databases. *IEEE Trans. Knowl. Data Eng.* 8, 4 (1996), 563–582.
- [24] Vitaly L. Khider and Grant E. Weddell. 2003. Reasoning about Uniqueness Constraints in Object Relational Databases. *IEEE Trans. Knowl. Data Eng.* 15, 5 (2003), 1295–1306.
- [25] Henning Köhler. 2006. Finding Faithful Boyce-Codd Normal Form Decompositions. In *AAIM*. 102–113.
- [26] Henning Köhler and Sebastian Link. 2016. SQL Schema Design: Foundations, Normal Forms, and Normalization. In *SIGMOD*. 267–279.
- [27] Nemanja Kojic and Dragan Milicev. 2020. Equilibrium of Redundancy in Relational Model for Optimized Data Retrieval. *IEEE Trans. Knowl. Data Eng.* 32, 9 (2020), 1707–1721.
- [28] Solmaz Kolahi and Leonid Libkin. 2010. An information-theoretic analysis of worst-case redundancy in database design. *ACM Trans. Database Syst.* 35, 1 (2010), 5:1–5:32.
- [29] Sebastian Kruse and Felix Naumann. 2018. Efficient Discovery of Approximate Dependencies. *Proc. VLDB Endow.* 11, 7 (2018), 759–772.
- [30] Jens Lechtenbörger and Gottfried Vossen. 2003. Multidimensional normal forms for warehouse design. *Inf. Syst.* 28, 5 (2003), 415–434.
- [31] Maurizio Lenzerini and Gaetano Santucci. 1983. Cardinality Constraints in the Entity-Relationship Model. In *ER*. 529–549.
- [32] Mark Levene and George Loizou. 2003. Why is the snowflake schema a good data warehouse design? *Inf. Syst.* 28, 3 (2003), 225–240.
- [33] Mark Levene and Millist W. Vincent. 2000. Justification for Inclusion Dependency Normal Form. *IEEE Trans. Knowl. Data Eng.* 12, 2 (2000), 281–291.
- [34] Stephen W. Liddle, David W. Embley, and Scott N. Woodfield. 1993. Cardinality Constraints in Semantic Data Models. *Data Knowl. Eng.* 11, 3 (1993), 235–270.
- [35] Sebastian Link and Henri Prade. 2019. Relational database schema design for uncertain data. *Inf. Syst.* 84 (2019), 88–110.
- [36] S. Link and Z. Wei. 2021. *Logical Schema Design that Quantify Update Inefficiency and Join Efficiency*. Report CDMTCS-552. The University of Auckland. <http://www.cs.auckland.ac.nz/research/groups/CDMTCS/>
- [37] Zehou Liu and Stratos Idreos. 2016. Main Memory Adaptive Denormalization. In *SIGMOD*. 2253–2254.
- [38] David Maier. 1983. *The Theory of Relational Databases*. Computer Science Press.
- [39] Wai Yin Mok, Yiu-Kai Ng, and David W. Embley. 1996. A Normal Form for Precisely Characterizing Redundancy in Nested Relations. *ACM Trans. Database Syst.* 21, 1 (1996), 77–106.
- [40] Antoni Olivé. 2007. Cardinality Constraints. In *Conceptual Modeling of Information Systems*. Springer, 83–102.
- [41] Thorsten Papenbrock, Jens Ehrlich, Jannik Marten, Tommy Neubert, Jan-Peer Rudolph, Martin Schönberg, Jakob Zwiener, and Felix Naumann. 2015. Functional Dependency Discovery: An Experimental Evaluation of Seven Algorithms. *PVLDB* 8, 10 (2015), 1082–1093.
- [42] Thorsten Papenbrock and Felix Naumann. 2016. A Hybrid Approach to Functional Dependency Discovery. In *SIGMOD*. 821–833.
- [43] Jean-Marc Petit, Farouk Toumani, Jean-François Boulicaut, and Jacques Kouloumdjian. 1996. Towards the Reverse Engineering of Denormalized Relational Databases. In *ICDE*. 218–227.
- [44] Tania Roblot, Miika Hannula, and Sebastian Link. 2018. Probabilistic Cardinality Constraints. *VLDB J.* 27, 6 (2018), 771–795.
- [45] Stefanie Scherzinger and Sebastian Sidortschuck. 2020. An Empirical Study on the Design and Evolution of NoSQL Database Schemas. In *ER*. 441–455.
- [46] Christian Soutou. 1998. Relational Database Reverse Engineering: Algorithms to Extract Cardinality Constraints. *Data Knowl. Eng.* 28, 2 (1998), 161–207.
- [47] Millist W. Vincent. 1999. Semantic Foundations of 4NF in Relational Database Design. *Acta Inf.* 36, 3 (1999), 173–213.
- [48] Ziheng Wei and Sebastian Link. 2019. Discovery and Ranking of Functional Dependencies. In *ICDE*. 1526–1537.
- [49] Jaejun Yoo, Ki-Hoon Lee, and Young-Ho Jeon. 2018. Migration from RDBMS to NoSQL Using Column-level Denormalization and Atomic Aggregates. *J. Inf. Sci. Eng.* 34, 1 (2018), 243–259.